

Phase-4_Project

October 28, 2022

0.0.1 Phase 4 Project

This project was a collaboration between * Benito Ywani * Patrick Arnold * Ahmad Samiee

1 Project Overview

In this project we measured text content sentiment in Twitter for two technology firms. Using natural language processing (NLP) methods, we generated models to make predictions for ternary (negative, neutral and positive) classification output.

2 Technical Summary

Our best model for Twitter sentiment prediction was Support Vector Classification which gave an overall accuracy of about 67%. The model under-predicted for all three classifications relative to actual for both firms. It missed predicting correctly negative sentiment the most followed by positive and then neutral. This may be because the dataset contained the fewest negative sentiment tweets. There were a moderate number of positive tweets and the most were neutral tweets. However, our model was able to correctly predict the overall trend for classification ratios where tweets with neutral sentiment had the highest ratio followed by positive and then negative. In addition, our model was able to correctly predict which of the two firms had a higher positive sentiment (this was also the case for neutral and negative sentiments).

3 Business Problem

Our stakeholder an equity investment firm is considering to purchase equity in a technology company and they have narrowed down their selection between two companies. As an additional analysis tool they like to take look at (potential customer) sentiment toward these two final selections. Therefore, they want a machine learning model that can predict sentiment in social media platform like Twitter and also want to know how accurately it can predict sentiment.

3.1 Master Dataset

The data comes from CrowdFlower via data.world which are tweets from Twitter collected during 2011 at South by Southwest Conference. It consists of about 9000 tweets and the tweets are labelled as negative, neutral or positive.

3.2 EDA and Feature Engineering

Created a column to identify whether the tweet was about Apple, Google or neither. Duplicated tweets, tweets for which there is no sentiment label and tweets that were neither about Apple or Google were dropped. Created another column to convert text sentiment into a numerical where 0 is for negative, 1 for neutral and 2 for positive sentiment. The cleaned data was then split into apple and google dataframes. These were then split into train (60% of data) and test (40%) dataframes where the three classification ratios for sentiment were split proportionally. The trained datasets for apple and google were combined into a single dataframe to be used for model training. The tweet text was then processed for natural language processing(NLP) as described in `nlp_doc_preparer` function.

```
[1]: # Load all the libraries

import pandas as pd
import numpy as np
import string
import re
import time

import matplotlib.pyplot as plt
import seaborn as sns

import nltk
from nltk.tokenize import RegexpTokenizer
from nltk.corpus import stopwords, wordnet
from nltk import pos_tag
from nltk.stem import WordNetLemmatizer
from nltk.probability import FreqDist

from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, recall_score, precision_score, \
    ↪confusion_matrix
from sklearn.metrics import make_scorer
from sklearn.metrics import plot_confusion_matrix
from sklearn.dummy import DummyClassifier
from imblearn.over_sampling import SMOTE

from nltk import word_tokenize
from keras.preprocessing.sequence import pad_sequences
```

```

from keras.layers import Input, Dense, LSTM, Embedding
from keras.layers import Dropout, Activation, Bidirectional, GlobalMaxPool1D
from keras.models import Sequential
from keras import initializers, regularizers, constraints, optimizers, layers
from keras.preprocessing import text, sequence
from keras import metrics

%matplotlib inline

```

```

[2]: # Read file
filename = 'data/judge-1377884607_tweet_product_company.csv'

sentiments_df = pd.read_csv(filename, encoding= 'unicode_escape')

```

```

[3]: # Data overview
sentiments_df.head()

```

```

[3]:
                                tweet_text \
0  .@wesley83 I have a 3G iPhone. After 3 hrs twe...
1  @jessedee Know about @fludapp ? Awesome iPad/i...
2  @swonderlin Can not wait for #iPad 2 also. The...
3  @sxsw I hope this year's festival isn't as cra...
4  @sxtxstate great stuff on Fri #SXSW: Marissa M...

emotion_in_tweet_is_directed_at \
0                                iPhone
1                iPad or iPhone App
2                                iPad
3                iPad or iPhone App
4                                Google

is_there_an_emotion_directed_at_a_brand_or_product
0                                Negative emotion
1                                Positive emotion
2                                Positive emotion
3                                Negative emotion
4                                Positive emotion

```

```

[6]: # Randomly check tweet text for a couple of rows

pd.set_option('display.max_colwidth', None)
indexes = np.random.randint(0, len(sentiments_df), 6)
sentiments_df['tweet_text'].loc[indexes]

```

```

[6]: 2998    It is never more apparent than at #sxsw how nice it would be if apple
      made stuff w/ removable batteries. #alwayshavingtoplugin #gsdm
      77

```

```

I worship @mention {link} #SXSW
1885
Win an iPad at SXSW via @mention #sxsw {link}
2182
Austinjs autocorrects to Sisyphus on the iPhone. Just sayin'. #sxsw
8237          #sxsw Just got a Samsung Focus (windows 7) phone at the android
dev meetup. We're excited to try it out and build something.
1770                                     #SXSW 2011: Novelty
of iPad news apps fades fast among digital delegates {link}
Name: tweet_text, dtype: object

```

```
[7]: # Check number of records, data types and which columns have nan
```

```
sentiments_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9093 entries, 0 to 9092
Data columns (total 3 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   tweet_text                                9092 non-null   object
1   emotion_in_tweet_is_directed_at          3291 non-null   object
2   is_there_an_emotion_directed_at_a_brand_or_product  9093 non-null   object
dtypes: object(3)
memory usage: 213.2+ KB

```

```
[8]: # Check values for column 'emotion_in_tweet_is_directed_at'
```

```
sentiments_df['emotion_in_tweet_is_directed_at'].value_counts()
```

```

[8]: iPad          946
     Apple         661
     iPad or iPhone App  470
     Google         430
     iPhone         297
     Other Google product or service  293
     Android App      81
     Android          78
     Other Apple product or service   35
     Name: emotion_in_tweet_is_directed_at, dtype: int64

```

```
[9]: # Confirm number of NaN in column 'emotion_in_tweet_is_directed_at'
```

```
sentiments_df['emotion_in_tweet_is_directed_at'].isna().sum()
```

```
[9]: 5802
```

```
[10]: # Check values for column 'is_there_an_emotion_directed_at_a_brand_or_product'

sentiments_df['is_there_an_emotion_directed_at_a_brand_or_product'].
↳value_counts()
```

```
[10]: No emotion toward brand or product      5389
Positive emotion                             2978
Negative emotion                             570
I can't tell                                 156
Name: is_there_an_emotion_directed_at_a_brand_or_product, dtype: int64
```

```
[11]: def find_company_name(text,emotion_at):

    '''
    Go through text ('tweet_text' column) and determine whether it is about
    ↳apple or google and
    return either 'apple' or 'google' according to search.
    If cannot tell company id from text, check emotion_at
    ↳('emotion_in_tweet_is_directed_at'
    column) for comapy id. If cannot find company id return 'cannot tell'.
    '''

    revised_emotion_at = str(emotion_at).lower()

    revised_text = str(text).lower()

    apple = ['apple','iphone','ipad']
    google = ['google','android']

    apple_count = 0
    google_count = 0

    for a in apple:
        apple_count += sum(1 for _ in re.finditer(r'\b%s\b' % re.escape(a),
        ↳revised_text))

    for g in google:
        google_count += sum(1 for _ in re.finditer(r'\b%s\b' % re.escape(g),
        ↳revised_text))

    if apple_count > google_count:
        return 'apple'
    elif google_count > apple_count:
        return 'google'
    elif revised_emotion_at != 'nan' and revised_emotion_at != 'default':
```

```

    for idx, item in enumerate(revised_emotion_at.split()):
        if item in apple:
            return 'apple'
        elif item in google:
            return 'google'
        else:
            if idx == len(revised_emotion_at.split()):
                return 'cannot tell'
    else:
        return 'cannot tell'

```

```

[12]: # Create a new column called 'company_name' and find whether the tweet is about
      ↪ apple,
      # google or cannot tell using function find_company_name

sentiments_df['company_name'] = sentiments_df.apply(
    lambda s: ↪
    ↪ find_company_name(s['tweet_text'], s['emotion_in_tweet_is_directed_at']), axis=1)

sentiments_df['company_name'].value_counts()

```

```

[12]: apple          5331
      google          2814
      cannot tell      948
      Name: company_name, dtype: int64

```

```

[13]: # Tweets for which cannot tell whether comment directed to apple or google the
      # sentiment is mostly neutral (over 95%)

cannot_tell = sentiments_df[sentiments_df['company_name'] == 'cannot tell']
cannot_tell['is_there_an_emotion_directed_at_a_brand_or_product'].value_counts()

```

```

[13]: No emotion toward brand or product      911
      Positive emotion                        23
      I can't tell                          10
      Negative emotion                        4
      Name: is_there_an_emotion_directed_at_a_brand_or_product, dtype: int64

```

```

[14]: # Drop any row where 'tweet_text' is nan
      sentiments_df.drop(sentiments_df[sentiments_df['tweet_text'].isna()].
      ↪ index, inplace=True)

      # Drop any row where 'is_there_an_emotion_directed_at_a_brand_or_product' says
      ↪ 'I can't tell'
      column_name = 'is_there_an_emotion_directed_at_a_brand_or_product'
      match = 'I can\'t tell'

```

```

sentiments_df.drop(sentiments_df[sentiments_df[column_name] == match].index,
    inplace=True)

# Drop any rows where column 'company_name' says 'cannot tell'
sentiments_df.drop(sentiments_df[sentiments_df['company_name'] == 'cannot_
    tell'].index,
    inplace=True)

# Drop any duplicated rows
sentiments_df.drop(sentiments_df[sentiments_df.duplicated()].index, inplace=True)

sentiments_df['company_name'].value_counts()

```

```

[14]: apple      5231
      google     2748
      Name: company_name, dtype: int64

```

```

[15]: # Create a function that returns 0 for negative sentiment, 1 for neutral
      # and 2 for positive sentiment

def convert_emotion_tonumber(emotion):
    if emotion == 'Negative emotion':
        return 0
    elif emotion == 'Positive emotion':
        return 2
    else:
        return 1 # for neutral emotion

```

```

[16]: # Create a column called 'sentiment' and pass value from
      # 'is_there_an_emotion_directed_at_a_brand_or_product' to function
      # convert_emotion_tonumber.

sentiments_df['sentiment'] = sentiments_df[
    'is_there_an_emotion_directed_at_a_brand_or_product'].map(
    lambda s: convert_emotion_tonumber(s))

sentiments_df['sentiment'].value_counts(normalize=True)

```

```

[16]: 1    0.559845
      2    0.369345
      0    0.070811
      Name: sentiment, dtype: float64

```

```

[17]: sentiments_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7979 entries, 0 to 9092
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_text                            7979 non-null   object
1   emotion_in_tweet_is_directed_at      3273 non-null   object
2   is_there_an_emotion_directed_at_a_brand_or_product  7979 non-null   object
3   company_name                          7979 non-null   object
4   sentiment                             7979 non-null   int64
dtypes: int64(1), object(4)
memory usage: 374.0+ KB

```

```

[18]: # Create an intermediate dataframe that only contains columns
      ↳ 'tweet_text', 'company_name'
      # sentiment

      sentiments2_df = sentiments_df[['tweet_text', 'company_name', 'sentiment']].copy()

```

```

[19]: sentiments2_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7979 entries, 0 to 9092
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   tweet_text      7979 non-null   object
1   company_name    7979 non-null   object
2   sentiment       7979 non-null   int64
dtypes: int64(1), object(2)
memory usage: 249.3+ KB

```

```

[20]: # Determine approximate total number of unique words in sentiments['tweet_text']

      tweet_words = sentiments_df['tweet_text'].map(word_tokenize).values
      total_vocabulary = set(word for tweet in tweet_words for word in tweet)
      print(f'Total number of unique words in tweet_words: {len(total_vocabulary)}')

```

Total number of unique words in tweet_words: 12131

```

[21]: # Split sentiments2_df into two data frames where apple_df contains all tweets
      ↳ about apple
      # google_df contains all tweets about google

      apple_df = sentiments2_df[sentiments2_df['company_name'] == 'apple'].copy()
      google_df = sentiments2_df[sentiments2_df['company_name'] == 'google'].copy()

```




```
[22]: apple_df['sentiment'].value_counts(normalize=True)
```

```
[22]: 1    0.516536
      2    0.404512
      0    0.078952
      Name: sentiment, dtype: float64
```

```
[23]: def func(pct, allvals):
      absolute = int(np.round(pct/100.*np.sum(allvals)))
      return "{:.0f}%\n".format(pct, absolute)


      a_size = list(apple_df['sentiment'].value_counts())
      a_labels = ['neutral', 'positive', 'negative']
      fig,axs = plt.subplots()
      axs.pie(a_size,labels=a_labels,autopct=lambda pct:
      ↪func(pct,a_size),colors=['grey','green','red']);
      axs.set_title('2011 SWSX Twitter Sentiment for Apple');
      fig.savefig('apple_sentiment')
```

2011 SWSX Twitter Sentiment for Apple



```
[24]: g_size = list(google_df['sentiment'].value_counts())
      g_labels = ['neutral', 'positive', 'negative']
      fig,axs = plt.subplots()
      axs.pie(g_size,labels=g_labels,autopct=lambda pct:
      ↪func(pct,g_size),colors=['grey','green','red']);
      axs.set_title('2011 SWSX Twitter Sentiment for Google');
      fig.savefig('google_sentiment');
```

2011 SWSX Twitter Sentiment for Google



```
[25]: def get_wordnet_pos(tag):  
    '''  
    Translate nltk POS to wordnet tags  
    '''  
  
    if tag.startswith('J'):  
        return wordnet.ADJ  
    elif tag.startswith('V'):  
        return wordnet.VERB  
    elif tag.startswith('N'):  
        return wordnet.NOUN  
    elif tag.startswith('R'):  
        return wordnet.ADV  
    else:  
        return wordnet.NOUN
```

```
[26]: def nlp_doc_preparer(doc,array='no'):  
    '''  
    - Customize nltk stop_words to include all punctuation marks, numbers and  
    → acronym 'szsw'  
    - Split text into words containing letters  
    - Make all letters lower case  
    - Use pos_tag to mark up the words for a particular part of a speech  
    - Use get_wordnet_pos to convert pos_tag to wordnet_pos  
    - Convert word to root word with WordNetLemmatizer  
    - Return processed document as a string of words
```

```

'''

custom_sw = stopwords.words('english')
punctuation = [c for c in list(string.punctuation)]
numbers = [n for n in range(0,10)]
custom_sw.extend(punctuation+numbers)
custom_sw.extend(['sxsw', 'SXSW', 'Sxsw'])

regex_token = RegexpTokenizer(r"([a-zA-Z]+(?:'[a-z]+)?)")
doc = regex_token.tokenize(doc)

doc = [word.lower() for word in doc]
doc = [word for word in doc if word not in custom_sw]

doc = pos_tag(doc)
doc = [(word[0], get_wordnet_pos(word[1])) for word in doc]

lemmatizer = WordNetLemmatizer()
doc = [lemmatizer.lemmatize(word[0], word[1]) for word in doc]

if array == 'no':
    return ' '.join(doc)
else:
    return doc

```

[28]: *# Show sentiment breakdown for Apple and Google*

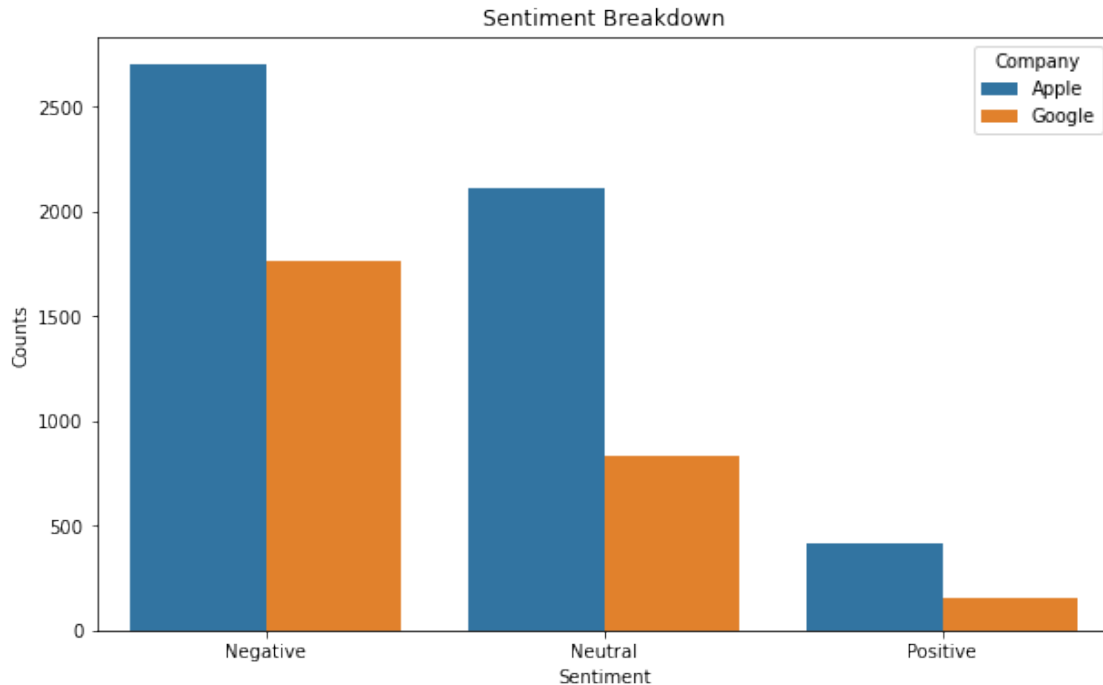
```

apple_sentiment = apple_df['sentiment'].value_counts()
google_sentiment = google_df['sentiment'].value_counts()

sentiments = pd.DataFrame({'Sentiment': ['Negative', 'Neutral', 'Positive',
                                           'Negative', 'Neutral', 'Positive'],
                           'Counts': (list(apple_sentiment) + list(google_sentiment)),
                           'Company': (['Apple']*3 + ['Google']*3)})

fig, ax = plt.subplots(figsize=(10,6))
sentiment_plot = sns.barplot(x='Sentiment', y='Counts', hue='Company',
    ↳data=sentiments);
sentiment_plot.set_title('Sentiment Breakdown');

```



[29]: *# Create data frames for each of the three sentiments for each company*

```
apple_neg_tw = apple_df[apple_df['sentiment'] == 0]
apple_neu_tw = apple_df[apple_df['sentiment'] == 1]
apple_pos_tw = apple_df[apple_df['sentiment'] == 2]

google_neg_tw = google_df[google_df['sentiment'] == 0]
google_neu_tw = google_df[google_df['sentiment'] == 1]
google_pos_tw = google_df[google_df['sentiment'] == 2]

app_neg_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪apple_neg_tw['tweet_text']]
app_neu_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪apple_neu_tw['tweet_text']]
app_pos_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪apple_pos_tw['tweet_text']]


goo_neg_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪google_neg_tw['tweet_text']]
goo_neu_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪google_neu_tw['tweet_text']]
goo_pos_tw = [nlp_doc_preparer(tweet, 'yes') for tweet in
    ↪google_pos_tw['tweet_text']]
```

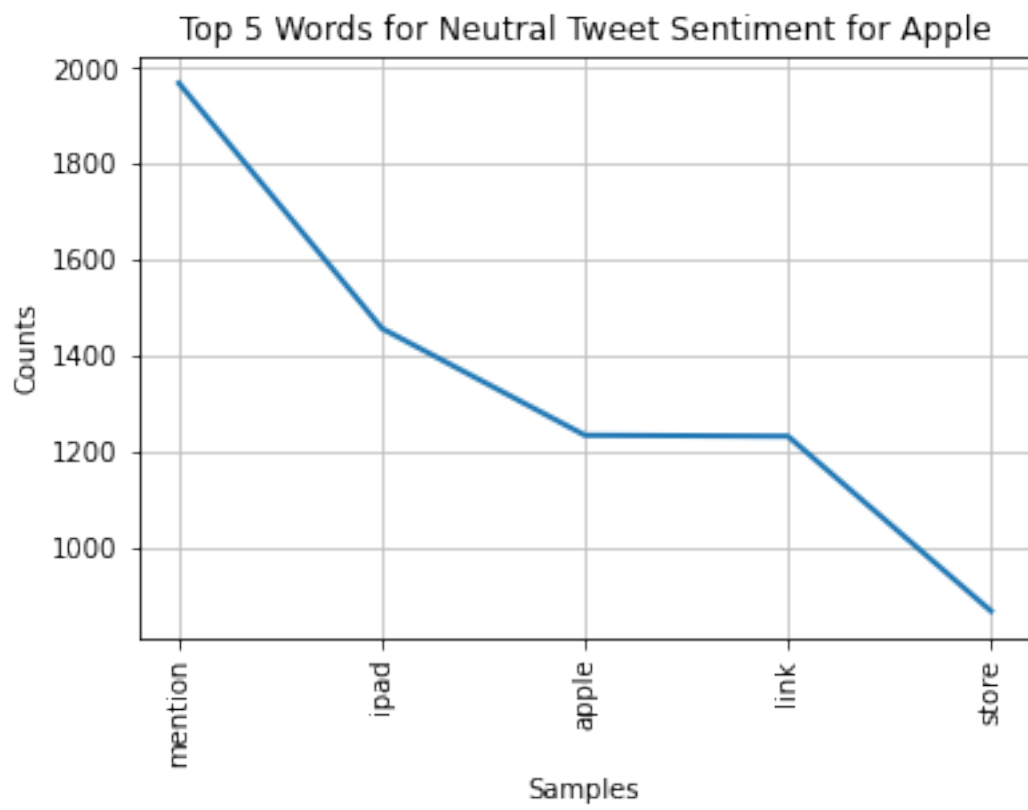
```
[30]: # A function to flatten array of arrays and return the flattened array
```


```
def flatten_list(list_of_lists):  
    return [word for line in list_of_lists for word in line]
```

```
[32]: # Show top 5 frequently occurring words for Apple for each sentiment
```

```
app_neg_tw_f = flatten_list(app_neg_tw)  
app_neu_tw_f = flatten_list(app_neu_tw)  
app_pos_tw_f = flatten_list(app_pos_tw)  
  
apple_tweets = [['Top 5 Words for Negative Tweet Sentiment for_  
↪Apple',app_neg_tw_f],  
                ['Top 5 Words for Neutral Tweet Sentiment for_  
↪Apple',app_neu_tw_f],  
                ['Top 5 Words for Positive Tweet Sentiment for_  
↪Apple',app_pos_tw_f]]  
  
fig, axs = plt.subplots()  
for i in range(3):  
    t = FreqDist(apple_tweets[i][1])  
    t.plot(5,title=apple_tweets[i][0]);
```








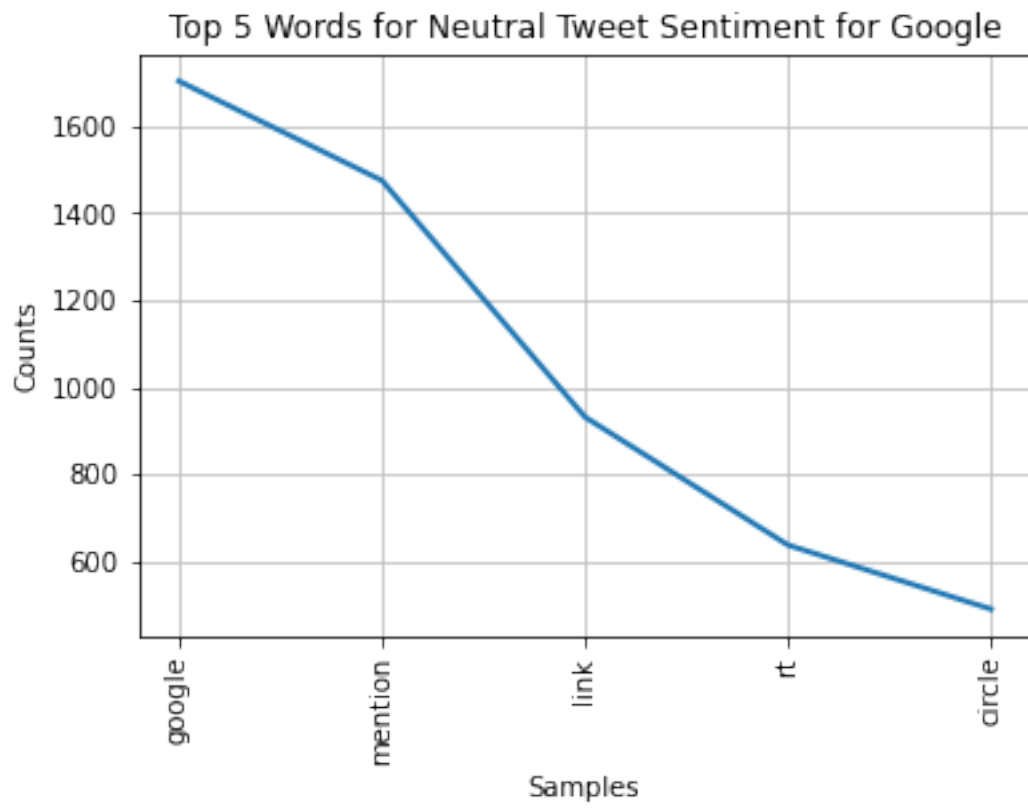
[33]: *# Show top 5 frequently occurring words for Google for each sentiment*

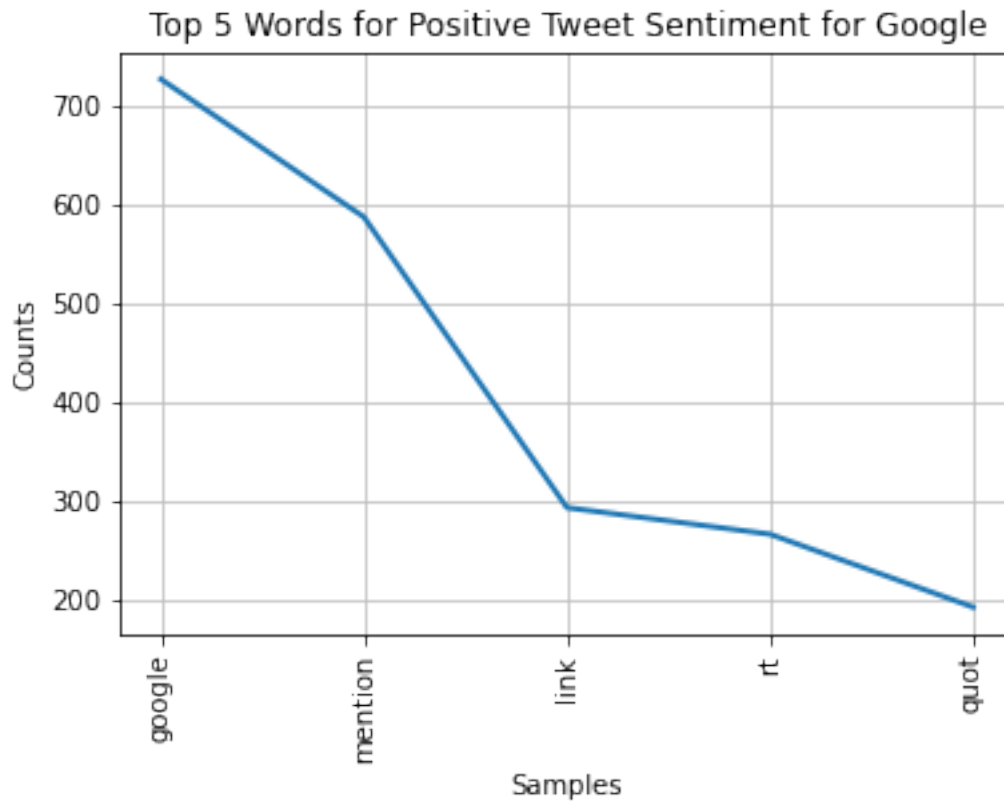
```
goo_neg_tw_f = flatten_list(goo_neg_tw)
goo_neu_tw_f = flatten_list(goo_neu_tw)
goo_pos_tw_f = flatten_list(goo_pos_tw)

apple_tweets = [['Top 5 Words for Negative Tweet Sentiment for_
↳Google',goo_neg_tw_f],
                ['Top 5 Words for Neutral Tweet Sentiment for_
↳Google',goo_neu_tw_f],
                ['Top 5 Words for Positive Tweet Sentiment for_
↳Google',goo_pos_tw_f]]

fig, axs = plt.subplots()
for i in range(3):
    t = FreqDist(apple_tweets[i][1])
    t.plot(5,title=apple_tweets[i][0]);
```







[]:

```
[34]: # Split apple_df and google_df into train and test samples using
      ↪ train_test_split

X_train_app, X_test_app, y_train_app, y_test_app =
      ↪ train_test_split(apple_df['tweet_text'],
                        apple_df['sentiment'], test_size=
      ↪ 0.4,
                        stratify =
      ↪ apple_df['sentiment'],
                        random_state = 6781)

X_train_goo, X_test_goo, y_train_goo, y_test_goo =
      ↪ train_test_split(google_df['tweet_text'],
                        google_df['sentiment'], test_size = 0.4,
                        stratify =
      ↪ google_df['sentiment'],
                        random_state = 6782)
```

```
[35]: a = len(X_train_app)
      g = len(X_train_goo)
      a+g
```

```
[35]: 4786
```

```
[ ]:
```

```
[36]: # Combined X_train_app and X_train_goo into single series
      # and y_train_app and y_train_goo

      X_train_combined = pd.concat([X_train_app,X_train_goo])
      y_train_combined = pd.concat([y_train_app,y_train_goo])
```

```
[37]: len(X_train_combined) == len(y_train_combined)
```

```
[37]: True
```

```
[38]: sum(X_train_combined.index == y_train_combined.index)
```

```
[38]: 4786
```

```
[39]: # Process X_train_combined, X_test_app, X_test_goo with nlp_doc_preparer_
      ↪function

      X_train_combined_processed = [nlp_doc_preparer(tweet) for tweet in_
      ↪X_train_combined]

      X_test_app_processed = [nlp_doc_preparer(tweet) for tweet in X_test_app]
      X_test_goo_processed = [nlp_doc_preparer(tweet) for tweet in X_test_goo]
```

```
[40]: def display_cross_validation_results(cross_val,model_name):
      '''
      Incorporate cross validation results into a pandas dataframe and display_
      ↪validation scores
      '''

      cross_val_results = pd.DataFrame(cross_val)

      accuracy = cross_val_results['mean_test_accuracy'][0]
      recall = cross_val_results['mean_test_recall'][0]
      precision = cross_val_results['mean_test_precision'][0]

      print('Validation','='*60)
      print(f'Validation results for {model_name}:')
      print(f'Accuracy: {accuracy}')
      print(f'Recall: {recall}')
```

```
print(f'Precision: {precision}\n')
```

```
[41]: def display_prediction_results(prediction, actual, company_name, model_name):  
    '''  
    Get predictions for X_test samples and display scores  
    '''  
  
    accuracy = accuracy_score(actual,prediction)  
    recall = recall_score(actual, prediction, average = 'micro')  
    precision = precision_score(actual, prediction, average = 'micro')  
  
    print('='*60)  
    print(f'Predictions for {company_name}: {model_name}')  
    print(f'Accuracy: {accuracy}')  
    print(f'Recall: {recall}')  
    print(f'Precision: {precision}\n')
```

```
[118]: def display_confusion_matrix_v2(actual_y,predicted_y,model_name,company_name):  
    conf_matrix = confusion_matrix(y_true=actual_y, y_pred=predicted_y)  
  
    fig, ax = plt.subplots(figsize=(6, 6))  
    ax.matshow(conf_matrix, cmap=plt.cm.Blues, alpha=0.3)  
  
    for i in range(conf_matrix.shape[0]):  
        for j in range(conf_matrix.shape[1]):  
            ax.text(x=j, y=i,s=conf_matrix[i, j], va='center', ha='center',  
↪size='xx-large')  
  
    plt.xlabel('Predictions', fontsize=14)  
    plt.ylabel('Actuals', fontsize=14)  
    plt.title(f'Confusion Matrix for {model_name}: {company_name}',  
↪fontsize=18);  
    save_image_as = 'confusion_matrix_' + '_' .join(model_name.split()) + ': ' +  
↪company_name  
    fig.savefig(save_image_as)
```

```
[42]: def display_confusion_matrix(model, X, y,company_name, model_name):  
    print(f'Confusion matrix for: {company_name}')  
    print(f'Model: {model_name}')  
    plot_confusion_matrix(model,X,y)  
    print('='*80)
```

3.3 Model Training and Model Prediction

Used CountVectorizer (and to lesser extent TfidfVectorizer) to convert text tweet into matrix of tokens. The models were trained and validated with the combined train dataframe from apple and google tweets. Used pipes to minimize redundant code and gridsearchcv for model tuning and validation. Model predictions were made with apple and google test dataframes.

```
[115]: # DummyClassifier
# Use pipes to countVectorize X_train_combined_processed and then train
# DummyClassifier

start = time.time()

cv_dm_pipe = Pipeline([('countvect',CountVectorizer()),
                        ('dm',DummyClassifier())])

grid_accuracy = make_scorer(accuracy_score)
grid_recall_micro = make_scorer(recall_score, average = 'micro')
grid_precision_micro = make_scorer(precision_score, average = 'micro')

cv_dm_params = {'countvect__input' : ['content'],
                'dm__random_state' : [1234],
                'dm__strategy' : ['prior']}

cv_dm_model_grid = GridSearchCV(estimator = cv_dm_pipe, param_grid =
                                cv_dm_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

cv_dm_model_grid.fit(X_train_combined_processed,y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
cv_dm_model_grid.best_params_
```

Training time: 0.4847831726074219

```
[115]: {'countvect__input': 'content',
        'dm__random_state': 1234,
        'dm__strategy': 'prior'}
```

```
[116]: # Display cross validation results, make prediction for X_test_app and
# X_test_goo
# diplay results. Use respective functions.

model_name = 'DummyClassifier with Count Vectorize'
```

```

display_cross_validation_results(cv_dm_model_grid.cv_results_,model_name)

y_test_app_hat = cv_dm_model_grid.predict(X_test_app_processed)
y_test_goo_hat = cv_dm_model_grid.predict(X_test_goo_processed)

display_prediction_results(y_test_app_hat, y_test_app, 'Apple',model_name)

display_prediction_results(y_test_goo_hat, y_test_goo, 'Google',model_name)

```

Validation =====

Validation results for DummyClassifier with Count Vectorize:

Accuracy: 0.5599666668848153

Recall: 0.5599666668848153

Precision: 0.5599666668848153

=====

Predictions for Apple: DummyClassifier with Count Vectorize

Accuracy: 0.5164835164835165

Recall: 0.5164835164835165

Precision: 0.5164835164835165

=====

Predictions for Google: DummyClassifier with Count Vectorize

Accuracy: 0.6418181818181818

Recall: 0.6418181818181818

Precision: 0.6418181818181818

```

[121]: # Display confusion matrix for X_test_app


#display_confusion_matrix_v2(y_test_app,y_test_app_hat,model_name,'Apple')
display_confusion_matrix(cv_dm_model_grid, X_test_app_processed,
                          y_test_app, 'Apple',model_name)

```

Confusion matrix for: Apple

Model: DummyClassifier with Count Vectorize

=====




```
[122]: # Display confusion matrix for X_test_goo
#display_confusion_matrix_v2(y_test_goo,y_test_goo_hat,model_name,'Google')
display_confusion_matrix(cv_dm_model_grid,X_test_goo_processed,
                          y_test_goo, 'Google',model_name)
```

Confusion matrix for: Google

Model: DummyClassifier with Count Vectorize

=====



```
[47]: # MultinomialNB
# Use pipes to countVectorize X_train_combined_processed and then train
↳ MultinomialNB

start = time.time()

cv_mnb_pipe = Pipeline([('countvect', CountVectorizer()),
                        ('mnb', MultinomialNB())])

cv_mnb_params = {'countvect__input' : ['content'], 'mnb__alpha' : [1.0]}

cv_mnb_model_grid = GridSearchCV(estimator = cv_mnb_pipe, param_grid =
↳ cv_mnb_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

cv_mnb_model_grid.fit(X_train_combined_processed, y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
```

Training time: 0.5058467388153076


```
[48]: # Display cross validation results, make prediction for X_test_app and
      ↪X_test_goo
      # display results. Use respective functions.

      model_name = 'MultinomialNB with Count Vectorize'

      display_cross_validation_results(cv_mnb_model_grid.cv_results_,model_name)

      y_test_app_hat2 = cv_mnb_model_grid.predict(X_test_app_processed)
      y_test_goo_hat2 = cv_mnb_model_grid.predict(X_test_goo_processed)

      display_prediction_results(y_test_app_hat2, y_test_app, 'Apple',model_name)

      display_prediction_results(y_test_goo_hat2, y_test_goo, 'Google',model_name)
```

```
Validation =====
Validation results for MultinomialNB with Count Vectorize:
Accuracy: 0.6276688852385346
Recall: 0.6276688852385346
Precision: 0.6276688852385346

=====
Predictions for Apple: MultinomialNB with Count Vectorize
Accuracy: 0.6086956521739131
Recall: 0.6086956521739131
Precision: 0.6086956521739131


=====
Predictions for Google: MultinomialNB with Count Vectorize
Accuracy: 0.6854545454545454
Recall: 0.6854545454545454
Precision: 0.6854545454545454
```

```
[49]: # Display confusion matrix for X_test_app

      display_confusion_matrix(cv_mnb_model_grid, X_test_app_processed,
                               y_test_app, 'Apple',model_name)
```

```
Confusion matrix for: Apple
Model: MultinomialNB with Count Vectorize
```

```
=====
```




```
[50]: # Display confusion matrix for X_test_goo
```

```
display_confusion_matrix(cv_mnb_model_grid,X_test_goo_processed,
                          y_test_goo, 'Google',model_name)
```

Confusion matrix for: Google

Model: MultinomialNB with Count Vectorize

=====



```
[51]: # Use pipes to TfidfVectorize X_train_combined_processed and then train
      ↪ MultinomialNB()

start = time.time()

tf_mnb_pipe = Pipeline([('tfvect', TfidfVectorizer()),
                        ('mnb', MultinomialNB())])

tf_mnb_params = {'tfvect__input' : ['content'], 'mnb__alpha' : [1.0]}

tf_mnb_model_grid = GridSearchCV(estimator = tf_mnb_pipe, param_grid =
      ↪ tf_mnb_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

tf_mnb_model_grid.fit(X_train_combined_processed, y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
```

Training time: 0.4806020259857178

```
[52]: # Display cross validation results, make prediction for X_test_app and
      ↪X_test_goo
      # display results. Use respective functions.

      model_name = 'MultinomialNB with TF-IDF Vectorize'

      display_cross_validation_results(tf_mnb_model_grid.cv_results_,model_name)

      y_test_app_hat3 = tf_mnb_model_grid.predict(X_test_app_processed)
      y_test_goo_hat3 = tf_mnb_model_grid.predict(X_test_goo_processed)

      display_prediction_results(y_test_app_hat3, y_test_app, 'Apple',model_name)

      display_prediction_results(y_test_goo_hat3, y_test_goo, 'Google',model_name)
```

Validation =====

Validation results for MultinomialNB with TF-IDF Vectorize:

Accuracy: 0.6305925135743002

Recall: 0.6305925135743002

Precision: 0.6305925135743002

=====

Predictions for Apple: MultinomialNB with TF-IDF Vectorize

Accuracy: 0.6077400860009555

Recall: 0.6077400860009555

Precision: 0.6077400860009555

=====

Predictions for Google: MultinomialNB with TF-IDF Vectorize

Accuracy: 0.6863636363636364

Recall: 0.6863636363636364

Precision: 0.6863636363636364


```
[53]: # Display confusion matrix for X_test_app

      display_confusion_matrix(tf_mnb_model_grid, X_test_app_processed,
                               y_test_app, 'Apple',model_name)
```

Confusion matrix for: Apple

Model: MultinomialNB with TF-IDF Vectorize

=====




```
[54]: # Display confusion matrix for X_test_goo
```

```
display_confusion_matrix(tf_mnb_model_grid,X_test_goo_processed,
                          y_test_goo, 'Google',model_name)
```

Confusion matrix for: Google

Model: MultinomialNB with TF-IDF Vectorize

=====



```
[55]: # KNeighborsClassifier
# Use pipes to CountVectorize X_train_combined_processed and then train
# → KNeighborsClassifier

start = time.time()

cv_knn_pipe = Pipeline([('countvect', CountVectorizer()),
                        ('knn', KNeighborsClassifier())])

cv_knn_params = {'countvect__input' : ['content'],
                 'knn__n_neighbors' : [7], # tried 3,7,21
                 'knn__weights' : ['distance'], # tried 'uniform'
                 'knn__leaf_size' : [7]} #tried 3,15,30,45

cv_knn_model_grid = GridSearchCV(estimator = cv_knn_pipe, param_grid =
    → cv_knn_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

cv_knn_model_grid.fit(X_train_combined_processed, y_train_combined)
end = time.time()
```

```
print(f'Training time: {end-start}')
cv_knn_model_grid.best_params_
```

Training time: 1.0253260135650635

```
[55]: {'countvect__input': 'content',
      'knn__leaf_size': 7,
      'knn__n_neighbors': 7,
      'knn__weights': 'distance'}
```

```
[56]: # Display cross validation results, make prediction for X_test_app and
      ↪X_test_goo
      # display results. Use respective functions.

      model_name = 'KNeighborsClassifier with Count Vectorize'

      display_cross_validation_results(cv_knn_model_grid.cv_results_,model_name)

      y_test_app_hat4 = cv_knn_model_grid.predict(X_test_app_processed)
      y_test_goo_hat4 = cv_knn_model_grid.predict(X_test_goo_processed)

      display_prediction_results(y_test_app_hat4, y_test_app, 'Apple',model_name)

      display_prediction_results(y_test_goo_hat4, y_test_goo, 'Google',model_name)
```

Validation =====

Validation results for KNeighborsClassifier with Count Vectorize:

Accuracy: 0.5927709897186537

Recall: 0.5927709897186537

Precision: 0.5927709897186537

=====

Predictions for Apple: KNeighborsClassifier with Count Vectorize

Accuracy: 0.5656951743908266

Recall: 0.5656951743908266

Precision: 0.5656951743908266

=====

Predictions for Google: KNeighborsClassifier with Count Vectorize

Accuracy: 0.6563636363636364

Recall: 0.6563636363636364

Precision: 0.6563636363636364

```
[57]: # Display confusion matrix for X_test_app


      display_confusion_matrix(cv_knn_model_grid, X_test_app_processed,
```

```
y_test_app, 'Apple',model_name)
```

Confusion matrix for: Apple

Model: KNeighborsClassifier with Count Vectorize

=====




```
[58]: # Display confusion matrix for X_test_goo
```

```
display_confusion_matrix(cv_knn_model_grid, X_test_goo_processed,  
                          y_test_goo, 'Google',model_name)
```

Confusion matrix for: Google

Model: KNeighborsClassifier with Count Vectorize

=====



```
[59]: # Use pipes to TfidfVectorize X_train_combined_processed and then train
      ↪ KNeighborsClassifier

start = time.time()

tf_knn_pipe = Pipeline([('tfvect', TfidfVectorizer()),
                        ('knn', KNeighborsClassifier())])

tf_knn_params = {'tfvect__input' : ['content'],
                  'knn__n_neighbors' : [7], # tried 3,7,21
                  'knn__weights' : ['distance'], # tried 'uniform'
                  'knn__leaf_size' : [3]} #tried 30,45

tf_knn_model_grid = GridSearchCV(estimator = cv_knn_pipe, param_grid =
      ↪ cv_knn_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

tf_knn_model_grid.fit(X_train_combined_processed, y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
```

```
tf_knn_model_grid.best_params_
```

Training time: 1.0332210063934326

```
[59]: {'countvect__input': 'content',  
      'knn__leaf_size': 7,  
      'knn__n_neighbors': 7,  
      'knn__weights': 'distance'}
```

```
[60]: # Display cross validation results, make prediction for X_test_app and  
      ↪X_test_goo  
      # diplay results. Use respective functions.  
  
model_name = 'KNeighborsClassifier with TF-IDF Vectorize'  
  
display_cross_validation_results(tf_knn_model_grid.cv_results_,model_name)  
  
y_test_app_hat5 = tf_knn_model_grid.predict(X_test_app)  
y_test_goo_hat5 = tf_knn_model_grid.predict(X_test_goo)  
  
display_prediction_results(y_test_app_hat5, y_test_app, 'Apple',model_name)  
  
display_prediction_results(y_test_goo_hat5, y_test_goo, 'Google',model_name)
```

```
Validation =====  
Validation results for KNeighborsClassifier with TF-IDF Vectorize:  
Accuracy: 0.5927709897186537  
Recall: 0.5927709897186537  
Precision: 0.5927709897186537
```

```
=====
```

```
Predictions for Apple: KNeighborsClassifier with TF-IDF Vectorize  
Accuracy: 0.5642618251313903  
Recall: 0.5642618251313903  
Precision: 0.5642618251313903
```


```
=====
```

```
Predictions for Google: KNeighborsClassifier with TF-IDF Vectorize  
Accuracy: 0.6327272727272727  
Recall: 0.6327272727272727  
Precision: 0.6327272727272727
```

```
[61]: # Diplay confusion matrix for X_test_app  
  
display_confusion_matrix(tf_knn_model_grid, X_test_app_processed,  
                          y_test_app, 'Apple',model_name)
```

Confusion matrix for: Apple
Model: KNeighborsClassifier with TF-IDF Vectorize


=====



```
[62]: # Display confusion matrix for X_test_app  
  
display_confusion_matrix(tf_knn_model_grid, X_test_goo_processed,  
                           y_test_goo, 'Google', model_name)
```

Confusion matrix for: Google
Model: KNeighborsClassifier with TF-IDF Vectorize

=====



[]:

```
[63]: # RandomForestClassifier
# Use pipes to CountVectorize X_train_combined_processed and then train
# → RandomForestClassifier

start = time.time()

cv_rf_pipe = Pipeline([('countvect', CountVectorizer()),
                        ('rf', RandomForestClassifier())])

cv_rf_params = {'countvect__input' : ['content'],
                'rf__random_state' : [42],
                'rf__max_depth' : [87], # tried 3,11,57,121
                'rf__criterion' : ['gini'], # log_loss gives warning
                'rf__max_features' : [None]} # tried 'sqrt' and 'log2'

cv_rf_model_grid = GridSearchCV(estimator = cv_rf_pipe, param_grid =
                                # → cv_rf_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')
```

```

cv_rf_model_grid.fit(X_train_combined_processed,y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
cv_rf_model_grid.best_params_

```

Training time: 105.48740887641907

```

[63]: {'countvect__input': 'content',
      'rf__criterion': 'gini',
      'rf__max_depth': 87,
      'rf__max_features': None,
      'rf__random_state': 42}

```

```

[64]: # Display cross validation results, make prediction for X_test_app and
      ↪ X_test_goo
      # display results. Use respective functions.

model_name = 'RandomForestClassifier with Count Vectorize'

display_cross_validation_results(cv_rf_model_grid.cv_results_,model_name)

y_test_app_hat6 = cv_rf_model_grid.predict(X_test_app_processed)
y_test_goo_hat6 = cv_rf_model_grid.predict(X_test_goo_processed)

display_prediction_results(y_test_app_hat6, y_test_app,'Apple',model_name)

display_prediction_results(y_test_goo_hat6, y_test_goo, 'Google',model_name)

```

```

Validation =====
Validation results for RandomForestClassifier with Count Vectorize:
Accuracy: 0.6385335610805339
Recall: 0.6385335610805339
Precision: 0.6385335610805339

```

```

=====
Predictions for Apple: RandomForestClassifier with Count Vectorize
Accuracy: 0.6144290492116579
Recall: 0.6144290492116579
Precision: 0.6144290492116579

```

```

=====
Predictions for Google: RandomForestClassifier with Count Vectorize
Accuracy: 0.7054545454545454
Recall: 0.7054545454545454
Precision: 0.7054545454545454

```


```
[65]: # Display confusion matrix for X_test_app

display_confusion_matrix(cv_rf_model_grid, X_test_app_processed,
                        y_test_app, 'Apple', model_name)
```

Confusion matrix for: Apple

Model: RandomForestClassifier with Count Vectorize

=====




```
[66]: # Display confusion matrix for X_test_app

display_confusion_matrix(cv_rf_model_grid, X_test_goo_processed,
                        y_test_goo, 'Google', model_name)
```

Confusion matrix for: Google

Model: RandomForestClassifier with Count Vectorize

=====



[]:

```
[67]: # LogisticRegression
# Use pipes to CountVectorize X_train_combined_processed and then train
↳ LogisticRegression

start = time.time()

cv_lr_pipe = Pipeline([('countvect', CountVectorizer()),
                        ('lr', LogisticRegression())])

cv_lr_params = {'countvect__input' : ['content'],
                'lr__random_state' : [321],
                'lr__penalty' : ['l2'],
                'lr__C' : [1.0], # tried 1 and 1000
                'lr__class_weight' : [None], # tried None and 'balanced'
                'lr__solver' : ['liblinear'], # tried 'newton-cg', 'lbfgs',
↳ 'liblinear', 'sag', 'saga'
                'lr__max_iter' : [1000]} # tried 100, 1000 and 10000

cv_lr_model_grid = GridSearchCV(estimator = cv_lr_pipe, param_grid =
↳ cv_lr_params,
                                scoring = {'accuracy' : grid_accuracy,
                                            'recall' : grid_recall_micro,
```

```

        'precision' : grid_precision_micro},
        refit = 'accuracy')

cv_lr_model_grid.fit(X_train_combined_processed,y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
cv_lr_model_grid.best_params_

```

Training time: 0.8365418910980225

```

[67]: {'countvect__input': 'content',
      'lr__C': 1.0,
      'lr__class_weight': None,
      'lr__max_iter': 1000,
      'lr__penalty': 'l2',
      'lr__random_state': 321,
      'lr__solver': 'liblinear'}

```

```

[68]: # Display cross validation results, make prediction for X_test_app and
      ↪X_test_goo
      # display results. Use respective functions.

model_name = 'LogisticRegression with Count Vectorize'

display_cross_validation_results(cv_lr_model_grid.cv_results_,model_name)

y_test_app_hat7 = cv_lr_model_grid.predict(X_test_app_processed)
y_test_goo_hat7 = cv_lr_model_grid.predict(X_test_goo_processed)

display_prediction_results(y_test_app_hat7, y_test_app,'Apple',model_name)

display_prediction_results(y_test_goo_hat7, y_test_goo, 'Google',model_name)

```

Validation =====

Validation results for LogisticRegression with Count Vectorize:

Accuracy: 0.6519049831698309

Recall: 0.6519049831698309

Precision: 0.6519049831698309

=====

Predictions for Apple: LogisticRegression with Count Vectorize

Accuracy: 0.630673674151935

Recall: 0.630673674151935


Precision: 0.630673674151935

=====

Predictions for Google: LogisticRegression with Count Vectorize
Accuracy: 0.7027272727272728
Recall: 0.7027272727272728
Precision: 0.7027272727272728


```
[69]: # Display confusion matrix for X_test_app  
  
display_confusion_matrix(cv_lr_model_grid, X_test_app_processed,  
                          y_test_app, 'Apple', model_name)
```

Confusion matrix for: Apple
Model: LogisticRegression with Count Vectorize
=====



```
[70]: # Display confusion matrix for X_test_app  
  
display_confusion_matrix(cv_lr_model_grid, X_test_goo_processed,  
                          y_test_goo, 'Google', model_name)
```

Confusion matrix for: Google
Model: LogisticRegression with Count Vectorize
=====



[]:

```
[123]: # SVC
# Use pipes to CountVectorize X_train_combined_processed and then train SVC

start = time.time()

cv_svc_pipe = Pipeline([('countvect', CountVectorizer()),
                        ('svc', SVC())])

cv_svc_params = {'countvect__input' : ['content'],
                  'svc__random_state': [234],
                  'svc__kernel' : ['rbf'], # tried 'poly', 'rbf', 'sigmoid'
                  'svc__degree' : [3], # tried 3,4,5
                  'svc__C' : [1.0], # tried 1 and 1000
                  'svc__class_weight' : [None]} # tried None and 'balanced'

cv_svc_model_grid = GridSearchCV(estimator = cv_svc_pipe, param_grid = [
    ↪cv_svc_params,
                                   scoring = {'accuracy' : grid_accuracy,
                                              'recall' : grid_recall_micro,
                                              'precision' : grid_precision_micro},
                                   refit = 'accuracy')
```

```

cv_svc_model_grid.fit(X_train_combined_processed,y_train_combined)
end = time.time()
print(f'Training time: {end-start}')
cv_svc_model_grid.best_params_

```

Training time: 15.133240222930908

```

[123]: {'countvect__input': 'content',
       'svc__C': 1.0,
       'svc__class_weight': None,
       'svc__degree': 3,
       'svc__kernel': 'rbf',
       'svc__random_state': 234}

```

```

[124]: # Display cross validation results, make prediction for X_test_app and
       ↪ X_test_goo
       # display results. Use respective functions.

model_name = 'SVC with Count Vectorize'

display_cross_validation_results(cv_svc_model_grid.cv_results_,model_name)

y_test_app_hat8 = cv_svc_model_grid.predict(X_test_app_processed)
y_test_goo_hat8 = cv_svc_model_grid.predict(X_test_goo_processed)

display_prediction_results(y_test_app_hat8, y_test_app,'Apple',model_name)

display_prediction_results(y_test_goo_hat8, y_test_goo, 'Google',model_name)

```

```

Validation =====
Validation results for SVC with Count Vectorize:
Accuracy: 0.6487671328503521
Recall: 0.6487671328503521
Precision: 0.6487671328503521

```

```

=====
Predictions for Apple: SVC with Count Vectorize
Accuracy: 0.6316292403248925
Recall: 0.6316292403248925
Precision: 0.6316292403248925

```


```

=====
Predictions for Google: SVC with Count Vectorize
Accuracy: 0.7127272727272728
Recall: 0.7127272727272728
Precision: 0.7127272727272728

```


```
[128]: # Display confusion matrix for X_test_app
#display_confusion_matrix_v2(y_test_app,y_test_app_hat8,model_name,'Apple')
display_confusion_matrix(cv_svc_model_grid, X_test_app_processed,
                        y_test_app, 'Apple',model_name)
```

Confusion matrix for: Apple
Model: SVC with Count Vectorize



```
[129]: # Display confusion matrix for X_test_app
#display_confusion_matrix_v2(y_test_goo,y_test_goo_hat8,model_name,'Google')
display_confusion_matrix(cv_svc_model_grid, X_test_goo_processed,
                        y_test_goo, 'Google',model_name)
```

Confusion matrix for: Google
Model: SVC with Count Vectorize



```
[ ]:
```

```
[84]: # Balance classification ratio with SMOTE for X_train_combined_processed
```

```
cvect = CountVectorizer()
X_train_combined_vected = cvect.fit_transform(X_train_combined_processed)

sm = SMOTE(random_state=3211)
X_train_combined_vec_sm, y_train_combined_sm = sm.fit_resample(
    X_train_combined_vected, y_train_combined)
```

```
[78]: # Classification ratio without SMOTE
```

```
y_train_combined.value_counts()
```

```
[78]: 1    2680
      2    1767
      0     339
      Name: sentiment, dtype: int64
```

```
[79]: # Classification ratio with SMOTE
```

```
y_train_combined_sm.value_counts()
```

```
[79]: 2    2680
      1    2680
      0    2680
      Name: sentiment, dtype: int64
```

```
[80]: # Train SVC model with SMOTED data

start = time.time()

sm_svc = SVC()

sm_svc_params = {'random_state' : [234], 'kernel' : ['rbf'], 'degree' : [3], 'C' : [1.0], 'class_weight' : [None]}

sm_svc_model_grid = GridSearchCV(estimator = sm_svc, param_grid = sm_svc_params,
                                scoring = {'accuracy' : grid_accuracy,
                                           'recall' : grid_recall_micro,
                                           'precision' : grid_precision_micro},
                                refit = 'accuracy')

sm_svc_model_grid.fit(X_train_combined_vec_sm, y_train_combined_sm)
end = time.time()
print(f'Training time: {end-start}')
sm_svc_model_grid.best_params_
```

Training time: 28.189348220825195

```
[80]: {'C': 1.0,
      'class_weight': None,
      'degree': 3,
      'kernel': 'rbf',
      'random_state': 234}
```

```
[81]: # Display cross validation results, make prediction for X_test_app and
      ↪ X_test_goo
      # display results. Use respective functions.

model_name = 'SVC with CountVectorizer and SMOTE'

X_test_app_vec = cvect.transform(X_test_app_processed)
X_test_goo_vec = cvect.transform(X_test_goo_processed)

display_cross_validation_results(sm_svc_model_grid.cv_results_, model_name)

y_test_app_hat9 = sm_svc_model_grid.predict(X_test_app_vec)
y_test_goo_hat9 = sm_svc_model_grid.predict(X_test_goo_vec)
```

```
display_prediction_results(y_test_app_hat9, y_test_app, 'Apple', model_name)

display_prediction_results(y_test_goo_hat9, y_test_goo, 'Google', model_name)
```

Validation =====

Validation results for SVC with CountVectorizer and SMOTE:

Accuracy: 0.6662935323383085

Recall: 0.6662935323383085

Precision: 0.6662935323383085

=====

Predictions for Apple: SVC with CountVectorizer and SMOTE

Accuracy: 0.5862398471094123

Recall: 0.5862398471094123

Precision: 0.5862398471094123

=====

Predictions for Google: SVC with CountVectorizer and SMOTE

Accuracy: 0.6581818181818182

Recall: 0.6581818181818182

Precision: 0.6581818181818182


[82]: *# Display confusion matrix for X_test_app*

```
display_confusion_matrix(sm_svc_model_grid, X_test_app_vec,
                          y_test_app, 'Apple', model_name)
```

Confusion matrix for: Apple

Model: SVC with CountVectorizer and SMOTE

=====




```
[85]: # Display confusion matrix for X_test_app

display_confusion_matrix(sm_svc_model_grid, X_test_goo_vec,
                          y_test_goo, 'Google', model_name)
```

Confusion matrix for: Google

Model: SVC with CountVectorizer and SMOTE

=====



[]:

[90]: *# Process target variable into one-hot encoded using pandas get_dummies*

```
y_train_comb_dum = pd.get_dummies(y_train_combined).values

y_test_app_dum = pd.get_dummies(y_test_app).values
y_test_goo_dum = pd.get_dummies(y_test_goo).values
```

[86]: *# Tokenize and format X_train_combined and X_test_app & X_test_goo features*

```
tokenizer = text.Tokenizer(num_words=10000)
tokenizer.fit_on_texts(list(X_train_combined))

X_train_comb_tokens = tokenizer.texts_to_sequences(X_train_combined)
X_train_comb_tok = sequence.pad_sequences(X_train_comb_tokens, maxlen=200)

X_test_app_tokens = tokenizer.texts_to_sequences(X_test_app)
X_test_app_tok = sequence.pad_sequences(X_test_app_tokens)

X_test_goo_tokens = tokenizer.texts_to_sequences(X_test_goo)
X_test_goo_tok = sequence.pad_sequences(X_test_goo_tokens)
```

```
[87]: # Instaniate Sequential model and add layers
```

```
sq_model = Sequential()
embedding_size = 4
sq_model.add(Embedding(10000, embedding_size))
sq_model.add(LSTM(25, return_sequences=True))
sq_model.add(GlobalMaxPool1D())
sq_model.add(Dropout(0.6))
sq_model.add(Dense(50, activation='relu'))
sq_model.add(Dropout(0.6))
sq_model.add(Dense(3, activation='softmax'))
```

```
[88]: # Specify compiler parameters
```

```
sq_model.compile(loss='categorical_crossentropy',
                 optimizer='adam',
                 metrics=['accuracy', 'Recall', 'Precision'])
```

```
[419]: sq_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 4)	40000
lstm_1 (LSTM)	(None, None, 25)	3000
global_max_pooling1d_1 (Glob	(None, 25)	0
dropout_2 (Dropout)	(None, 25)	0
dense_2 (Dense)	(None, 50)	1300
dropout_3 (Dropout)	(None, 50)	0
dense_3 (Dense)	(None, 3)	153

Total params: 44,453
Trainable params: 44,453
Non-trainable params: 0

```
[91]: # Train model with X_train_comb and y_train_comb_dum
```

```
sq_model.fit(X_train_comb_tok, y_train_comb_dum, epochs=5,
             batch_size=16, validation_split=0.1)
```

```
Epoch 1/5
270/270 [=====] - 21s 77ms/step - loss: 0.9424 -
accuracy: 0.5289 - recall: 0.2540 - precision: 0.5368 - val_loss: 0.7991 -
val_accuracy: 0.6660 - val_recall: 0.6660 - val_precision: 0.6660
Epoch 2/5
270/270 [=====] - 20s 75ms/step - loss: 0.9017 -
accuracy: 0.5447 - recall: 0.3894 - precision: 0.5571 - val_loss: 0.7620 -
val_accuracy: 0.6660 - val_recall: 0.6660 - val_precision: 0.6660
Epoch 3/5
270/270 [=====] - 20s 74ms/step - loss: 0.8606 -
accuracy: 0.5519 - recall: 0.4202 - precision: 0.6056 - val_loss: 0.7038 -
val_accuracy: 0.6660 - val_recall: 0.6159 - val_precision: 0.7024
Epoch 4/5
270/270 [=====] - 21s 77ms/step - loss: 0.7985 -
accuracy: 0.5925 - recall: 0.4091 - precision: 0.7012 - val_loss: 0.7350 -
val_accuracy: 0.6889 - val_recall: 0.5908 - val_precision: 0.7128
Epoch 5/5
270/270 [=====] - 20s 74ms/step - loss: 0.7219 -
accuracy: 0.6752 - recall: 0.4876 - precision: 0.7468 - val_loss: 0.7489 -
val_accuracy: 0.6931 - val_recall: 0.6075 - val_precision: 0.7239
```

```
[91]: <tensorflow.python.keras.callbacks.History at 0x7f8d8c5eb430>
```

```
[92]: # Obtain model predictions for X_test_app and X_test_goo
```

```
y_test_app_pp = sq_model.predict(X_test_app_tok)
y_test_goo_pp = sq_model.predict(X_test_goo_tok)
```

```
[93]: # Convert probabilities to predicted sample target value
```

```
y_test_app_hat10 = np.argmax(y_test_app_pp, axis=-1)
y_test_goo_hat10 = np.argmax(y_test_goo_pp, axis=-1)
```

```
[94]: # Display prediction results for X_test_app and X_test_goo
```

```
display_prediction_results(y_test_app_hat10, y_test_app, 'Apple', 'Sequential')

display_prediction_results(y_test_goo_hat10, y_test_goo, 'Google', 'Sequential')
```

```
=====
Predictions for Apple: Sequential
Accuracy: 0.5590062111801242
Recall: 0.5590062111801242
Precision: 0.5590062111801242

=====
Predictions for Google: Sequential
Accuracy: 0.6463636363636364
```

Recall: 0.6463636363636364
Precision: 0.6463636363636364

3.4 Conclusion

The base model, DummyClassifier, had a validation accuracy of about 56% on the combined apple and google tweet dataframe (train dataframe). Predictions for apple and google tweet test samples were performed separately to be able to compare the two between one another for sentiment. The base model gave a prediction of 52% accuracy for apple and 64% for google with an overall average accuracy of 58%. Five additional models (Multinomial NB, KNeighbor, Random Forest, Logistic Regression and SupportVectorClassifier(SVC)) were then evaluated for their prediction performance and they all gave similar accuracy within 5% of each other. SVC model was selected as our best model as it gave overall better performance relative to the other four. Validation accuracy for SVC model was about 65%. For test samples, accuracy for apple was about 63% and for google 71% with an overall accuracy of 67% for our best model. Utilized gridsearchcv to adjust hyperparameters(tunning) for these models. Following NLP processing the tweet text was converted into matrix of tokens using mostly CountVectorizer for TfidfVectorizer diminished model performance a little bit. To correct for data classification imbalance (7% negative, 56% neutral and 37% positive sentiment) the combined trained data was synthetically balanced using SMOTE. However, this did not improve model performance. We also tried neural network Sequential model but were unable to increase accuracy.

Our best model for Twitter sentiment prediction was therefore SVC which gave an overall accuracy of about 67%. The model under-predicted for all three classifications relative to actual for both firms. Also, it missed predicting correctly negative sentiment the most followed by positive and then neutral. This may be because the dataset contained the fewest number of negative sentiment tweets. The dataset has a moderate number of positive tweets and the most for neutral tweets. However, our model was able to correctly predict the overall trend for classification ratios where tweets with neutral sentiment had the highest ratio followed by positive and then negative. In addition, our model was able to correctly predict which of the two firms had a higher positive sentiment.

3.5 Next Steps:

To potentially improve models performance * Take a closer look at tweets that had negative sentiment. Our best model missed these more than positive and neutral sentiments * Setup a model based on scoring words or phrases for negative, neutral or positive sentiment * Consider data from other social media platforms like review sites for products and services

[]: